Evidence-claim verification

Judge only whether this evidence claim is supported by this source record. resource is the original record retrieved using evidence.source_id.

Check each drug name, dose, test value, time, and other fact in statement. Set valid=true only if the record supports the entire claim; set valid=false if contradicted or undocumented.

Do not substitute other records or clinical knowledge for this source. In reason, name the source_id and the supporting or conflicting facts.

Whole-task verification

Verify that the question and reference answer are supported by the original data in evidence_records. Records have been retrieved from the environment, and evidence claims checked individually.

Do not assume patient facts absent from the records. Post-decision data may support the reference answer, but actual subsequent treatment is not the only correct option for a clinical decision task.

Return valid=false for an ambiguous question, insufficient evidence, or an incorrect reference answer, and explain the basis.

Reasoning-depth assessment

Inspect the question, evidence, and candidate answer. Score the simplest sufficient path that correctly meets the core requirements. A candidate answer is neither automatically correct nor the only solution path.

1: Fact retrieval. Read, match, count, summarize, or restate recorded diagnoses. Repeated records do not raise the level.

2: Simple judgment. Facts plus one clinical rule suffice. Multiple independent simple judgments do not imply level 3.

3: Necessary integration. The conclusion depends on at least two independent, indispensable facts and their relationship. Removing a key fact changes or invalidates the conclusion.

4: Decision trade-off. At least two reasonable options, conflicting benefits, risks, or constraints, and patient evidence that determines the choice.

Check shortcuts: does the question state the conclusion, a diagnosis label contain the answer, one record suffice, or generic advice suffice? More evidence, longer time spans, or longer answers alone add no points.

In reason, state the solution path, necessary facts and relationships, shortcuts, and level justification. For level 4, specify options and trade-offs. Do not fabricate unconfirmed patient facts.

Score aggregation and HO tasks

GPT-4o-mini and GPT-5.4-nano score independently. Take the median of two valid scores and round down. If either is missing, no aggregate level is formed.

Verified tasks with aggregate level >=3 are higher-order (HO) tasks.