Evidence Grounding and Strategy Evolution for Agent Task Generation from Longitudinal EHRs

Anonymous ACL submission

Abstract

Generating high-quality tasks for medical agents from longitudinal EHRs is challenging because patient records differ in the clinical relationships they support, while available tools restrict how evidence can be retrieved. Here, we propose MedEvolve, a framework that improves medical task generation at two complementary levels, task grounding and strategy evolution. At the task-construction level, dual exploration distills recurring clinical-temporal patterns and tool-access constraints into reusable guidance. At the strategy level, feedback-driven evolution then converts verification and independent-solving outcomes into cross-patient construction rules and task-direction preferences, rather than only repairing individual questions. MedEvolve achieves higher verified higher-order task yield than the baselines on both benchmarks. Adding its synthetic tasks improves official-test success rates by 6 and 2 percentage points for ACE and Reflexion, respectively.

1 Introduction

Longitudinal electronic health records (EHRs) document diagnoses, tests, and treatments across encounters and provide the basis for medical agent tasks. MedAgentBench (Jiang et al., 2025) and PhysicianBench (Liu et al., 2026) evaluate LLM agents in these environments, but expert task design and review limit scalability. TaskCraft (Shi et al., 2025) generates tasks from web pages, PDFs, and images through depth and width expansion; CuES (Mai et al., 2025) synthesizes tasks through environment exploration. Neither is designed for longitudinal EHR structure and evidence retrieval through read-only tools.

Task generation faces two challenges. First, task validity depends on both patient records and tool capabilities. A single trajectory does not reveal recurring clinical relationships or evidence distributions across patients, and complete trajectories may exceed the input budget. Whether the required evidence is retrievable must also be checked through tool interaction. Second, fixed guidance yields uneven task quality and difficulty across patients; defects often emerge only during verification or independent solving.

We propose MedEvolve, which combines data and environment exploration with feedback-driven strategy updates across patients, without changing model parameters. Our contributions are: (1) grounding task generation in clinical event patterns and tool-access constraints; (2) using evaluation feedback to revise natural-language guidance and task-direction weights; and (3) evaluating valid higher-order task yield in two EHR environments and the downstream value of synthetic tasks for experience learning.

2 Method

2.1 Problem Setting

Given a medical environment $\mathcal{E}$ and a set of longitudinal patient trajectories $\mathcal{D} = \{X_p\}$, the goal is to construct a medical agent task for each trajectory X_p. Each task consists of an instruction q, a candidate reference answer a, and source evidence e. Generation is controlled by a strategy $\theta_r = (g_r, w_r)$, where g_r is natural-language generation guidance, w_r holds sampling weights over task types and difficulty directions, and r is the evolution round. MedEvolve first obtains data guidance K_D and environment guidance K_E through exploration, then updates θ_r using evaluation feedback. Figure 1 shows the overall pipeline.

2.2 Data and Environment Exploration

A single patient's trajectory shows only which records that case contains, which makes it hard to tell what pattern those records follow across the whole data source, and reading a large number of patients again at every generation step is infeasible.

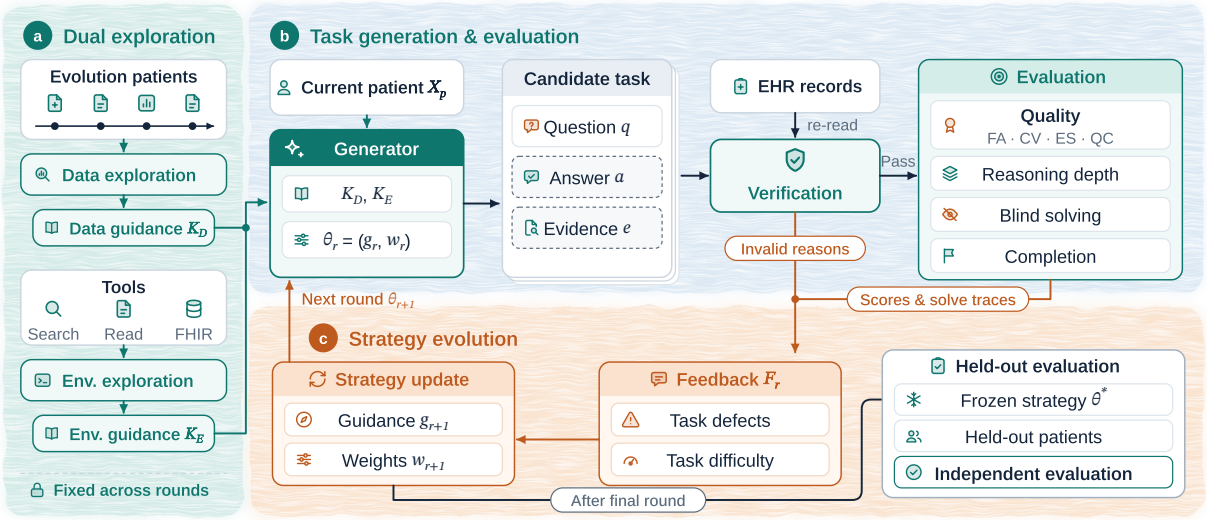


Figure 1: Overview of MedEvolve. (a) Data and environment exploration produce fixed guidance K_D and K_E. (b) Tasks generated from the patient trajectory, exploration guidance, and current strategy undergo source verification, quality and reasoning assessment, independent solving, and completion assessment. Reference answers and evidence claims are withheld from the solver. (c) Evaluation feedback updates generation guidance and direction weights; the final strategy is frozen for held-out evaluation.

We therefore run data exploration before generation and organize cross-patient observations into reusable guidance.

Data exploration targets longitudinal patient trajectories. The exploration agent reads the records of multiple patients, analyzes how the data are organized, how they change over time, how clinical events relate to one another, and how evidence is distributed, and distills the information patterns required by different task types into data guidance K_D. K_D describes only reusable patterns and contains no patient-specific clinical facts.

Environment exploration targets the environment interface. The exploration agent actually calls read-only query tools and records access methods, return structures, and interaction constraints into environment guidance K_E, which serves to judge whether the evidence a question requires can be obtained through the tools.

Both explorations are completed during initialization and remain unchanged in all subsequent rounds.

2.3 Task Generation and Evaluation

For patient X_p, the generator conditions on K_D, K_E, and the current strategy θ_r, and generates (q, a, e) according to the specified task type and difficulty direction, where key facts carry references to the original records.

Candidate tasks then pass through verification, scoring, and independent solving in sequence. Verification re-reads the original records and checks that the evidence claims support the reference answer; tasks that pass verification are then assessed for task quality and reasoning depth and answered independently by a fixed agent in the environment. The solving agent has no access to reference answers or evidence claims.

2.4 Feedback-Driven Strategy Updates

After each round, the system aggregates that round's verification defects, quality and reasoning scores, and solve traces, and attributes them into feedback F_r. Attribution must separate defects of the task itself (such as factual errors, insufficient evidence, or ambiguous questions), the genuine reasoning difficulty the task poses, and model, tool, or environment failures, so that execution anomalies are not mistaken for task difficulty. The strategy is updated as

$$\theta_{r+1} = \mathcal{U}(\theta_r, F_r), \tag{1}$$

where $\mathcal{U}$ is implemented by an LLM. The update acts on two parts: it writes generalizable construction experience into g_{r+1}, adjusting the selection of clinical information, the organization of evidence, and task design; and it adjusts w_{r+1} according to the quality and difficulty performance of each direction, changing the direction distribution of subsequent tasks.

Different rounds use different patient batches to generate tasks. After the scheduled rounds the strat-

Method	PhysicianBench		MedAgentBench	
	HO count (yield)	Mean level	HO count (yield)	Mean level
MedEvolve	**23 (14.2%)**	2.27	**77 (51.3%)**	**2.73**
EnvExplore	21 (13.0%)	2.25	64 (42.7%)	2.56
TaskCraft	0 (0.0%)	1.01	0 (0.0%)	1.00
Workflow	3 (1.9%)	**2.30**	2 (1.3%)	2.67
SingleShot	3 (1.9%)	1.20	11 (7.3%)	1.44
SelfRefine	4 (2.5%)	2.20	2 (1.3%)	2.00

Table 1: Valid higher-order task generation with GPT-5.4-nano. Higher-order (HO) tasks pass evidence verification and have an aggregate level of at least 3. Yield, shown in parentheses, uses the fixed number of requested tasks as its denominator (162 for PhysicianBench; 150 for MedAgentBench). Mean level covers all verified tasks with difficulty scores. Bold indicates the highest value in each column.

egy is frozen and evaluated on a held-out patient set.

3 Experiments

3.1 Evaluation Conditions and Metrics

Environments and comparison conditions. We use 108 PhysicianBench patients and 100 MedAgentBench patients with clinical records, reserving 81 and 75, respectively, for held-out evaluation and the remainder for strategy evolution. The sets are patient-disjoint. All methods share the held-out patients and generation model, GPT-5.4-nano, and generate at most two read-only tasks per patient, yielding fixed totals of 162 and 150 requested tasks. Patient splits and information-access rules appear in Appendix B.

Metrics. Tasks are verified against source records. GPT-4o-mini and GPT-5.4-nano independently assess the minimum reasoning required, with scores aggregated into levels 1–4. Verified tasks at level 3 or above count as higher-order (HO). Our primary metrics are HO count and yield, the latter dividing HO count by the fixed requested-task total and thus accounting for generation failures and invalid tasks. Mean level across all verified tasks with difficulty scores is an auxiliary metric. Appendix D details scoring and aggregation.

3.2 Experiment 1: Generating Valid Higher-Order Tasks

We compare generation methods under a shared patient set and task-count limit. SingleShot generates in one pass; Workflow uses localization, generation, and filtering; SelfRefine (Madaan et al., 2023) revises candidates through self-feedback; TaskCraft applies depth and width expansion. EnvExplore uses environment guidance with a fixed initial strategy, whereas MedEvolve combines data and environment exploration with the full strategy after two feedback-update rounds. Computational costs differ across methods; the comparison controls patients and task counts. Appendix C describes procedures and budgets.

MedEvolve achieves the highest HO yield in both environments (Table 1), with a larger margin over the closest baseline, EnvExplore, on MedAgentBench: 13 additional tasks and an 8.6-percentage-point gain, compared with two tasks and 1.2 points on PhysicianBench. The full framework therefore leads in this evaluation, although the PhysicianBench margin is small.

Verification rate and HO yield capture different task properties. TaskCraft produces no HO tasks and has a mean level near 1. MedEvolve has a lower verification rate than some baselines but retains the most valid HO tasks (Appendix F); verification rate alone does not capture reasoning demands. Mean level also requires attention to sample size: Workflow slightly exceeds MedEvolve on PhysicianBench (2.30 vs. 2.27), based on only 10 verified tasks.

3.3 Experiment 2: Supplementing Official Training Tasks with Synthetic Tasks

We test whether synthetic tasks supplement official tasks for experience learning. In PhysicianBench, ACE, Reflexion (Shinn et al., 2023), and OpenClaw each learn under two settings: *Official 20* uses 20 official tasks; *Official + Synthetic 40* adds 20 verified MedEvolve tasks to those same tasks. Both

settings use Qwen3.8-27B with fixed model parameters. Experience is frozen after training, and both settings are evaluated on the same 50 official test tasks. Reference answers are used only for scoring. Appendix E details framework mechanisms and reproducibility settings.

Framework	Official 20	Official + Synthetic 40	Δ (pp)
ACE	4.0% (2/50)	10.0% (5/50)	+6.0
Reflexion	12.0% (6/50)	14.0% (7/50)	+2.0
OpenClaw	24.0% (12/50)	20.0% (10/50)	−4.0

Table 2: Task success rates before and after adding synthetic tasks to official training tasks. Success requires the entire task to pass; all rates use the 50 planned test tasks as the denominator. Differences are in percentage points (pp). ACE has 48 scored tasks in the official-only setting and 50 in the mixed setting; all other settings have 50.

Adding synthetic tasks raises success rates by 6.0 percentage points for ACE and 2.0 for Reflexion, but lowers OpenClaw's rate by 4.0 points (Table 2). Synthetic tasks thus show potential in some experience-learning frameworks, with no consistent benefit across the evaluated frameworks. Their different experience representations and update mechanisms may affect task use, but this comparison does not establish the cause of the observed differences.

4 Conclusion

We present MedEvolve, which combines data and environment exploration to ground medical agent task generation in longitudinal patient records and tool-based evidence access. Verification and independent-solving feedback update generation guidance and task-direction weights across patients, without updating model parameters.

On PhysicianBench and MedAgentBench, MedEvolve achieves the highest valid higher-order task yield among the six compared methods, producing 23 and 77 verified higher-order tasks, respectively. Downstream experience-learning experiments further suggest potential value for synthetic tasks in some frameworks: adding synthetic tasks improves official-test success rates for ACE and Reflexion, but reduces the rate for OpenClaw, indicating differences in benefits across frameworks. These findings support dual exploration as an approach to constructing higher-order medical agent tasks and provide an empirical basis for studying how synthetic tasks interact with experience-learning mechanisms.

Limitations

The downstream experience-learning experiment compares training on 20 official tasks with training on those tasks plus 20 synthetic tasks, changing both task count and composition. Controls using an equal number of official tasks or tasks produced by other generation methods would help distinguish the effects of training-set size from those of synthetic-task characteristics. The current downstream results come from a single seed and 50 test tasks in PhysicianBench; the stability of the gains and their performance in other environments require further evaluation.

MedEvolve derives generation guidance from patient records and tool interfaces within a particular environment. The current held-out patient evaluation examines strategy transfer within that environment. The applicability of existing guidance, and the adaptations it requires when healthcare institutions, record structures, or evidence-retrieval tools change, have not yet been systematically evaluated.

References

Yixing Jiang et al. 2025. MedAgentBench: A realistic virtual EHR environment to benchmark medical LLM agents. *arXiv preprint arXiv:2501.14654*.

Timothy Kassis, Vinayak Agarwal, Yuhuan He, Darshil Patel, and Aubrey M. Brueckner. 2026. Scientific agent skills: A library of procedural knowledge for research agents. *arXiv preprint arXiv:2609.00065*.

Ruoqi Liu et al. 2026. PhysicianBench: Evaluating LLM agents in real-world EHR environments. *arXiv preprint arXiv:2605.02240*.

Aman Madaan et al. 2023. Self-Refine: Iterative refinement with self-feedback. *arXiv preprint arXiv:2303.17651*.

Shinji Mai et al. 2025. CuES: A curiosity-driven and environment-grounded synthesis framework for agentic RL. *arXiv preprint arXiv:2512.01311*.

Dingfeng Shi et al. 2025. TaskCraft: Automated generation of agentic tasks. *arXiv preprint arXiv:2506.10055*.

Noah Shinn et al. 2023. Reflexion: Language agents with verbal reinforcement learning. *arXiv preprint arXiv:2303.11366*.

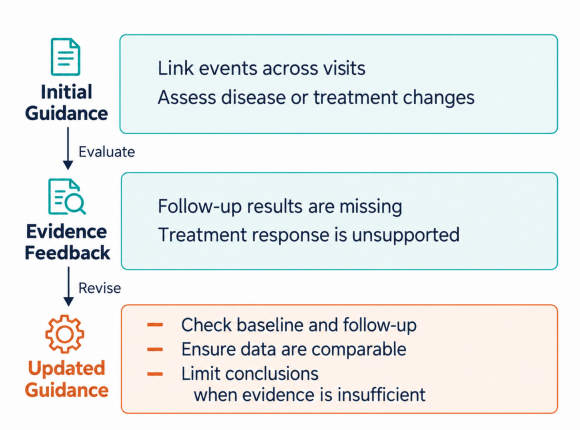


Figure 2: Example feedback-driven revision of generation guidance g_r. Verification feedback flags insufficient follow-up evidence; the revised guidance adds baseline/follow-up and comparability checks and restricts conclusions when evidence is weak. The example illustrates how g_r can change, not that updates improve held-out metrics. Text is paraphrased from one evolution trace.

A Exploration and Generation Interfaces

A.1 Data Exploration

Exploration tools expose open patient IDs, budgeted trajectories, and optional compressed summaries. Inputs retain patient identifiers, source pointers, and coverage flags; the output is patient-agnostic natural-language guidance K_D. The guide summarizes data organization, temporal change, event relations, evidence citation patterns, and supportable questioning conditions, without retaining individual clinical values or source IDs.

Compressed summaries locate candidate patterns; original records support concrete claims. Missing records and budget truncation are stated separately. Time windows, record counts, and clinical relations in K_D are inductive questioning constraints, not hard-coded medical rules applied uniformly across patients.

A.2 Environment Exploration

Environment exploration queries native read-only tools in an isolated instance over open patients. Guidance K_E records validated parameters, return shapes, pagination and truncation behavior, and typical retrieval procedures, and distinguishes call failures from unverified capabilities. Patient-specific observations stay out of the shared guide.

Table 3: Generation-side patient splits (seed 42).

Env	Pool	Evo.	Held-out	Rounds
PhysicianBench	108	27	81	13/14
MedAgentBench	100	25	75	12/13

A.3 Guidance Revision Example

A.4 Task Generation and Strategies

A task records patient ID, clinical category, question, candidate reference answer, source identifiers with per-source factual claims, and execution requirements. The shared generation interface leaves the temporal-view field empty and requires read-only execution. One generation call receives the budgeted trajectory, K_D, K_E, current guidance g_r, and two sampled directions $Z_{p,r}$, and returns at most two candidates.

The strategy is $\theta_r = (g_r, w_r)$ with natural-language guidance and direction weights over five clinical categories and difficulty descriptions. Directions are sampled proportional to weights (at least one positive). K_D and K_E stay fixed during updates; only g_r and w_r change. Evolution rounds sample from the current weights; held-out evaluation uses the frozen updated guidance, task-direction descriptions, and weights.

B Environments, Patients, and Information Access

Table 3 freezes evolution and held-out lists under seed 42 with prior exploration-exposure constraints; patients do not cross sets. Trajectories are ordered by record time and budgeted for model input; verification rereads original sources rather than input summaries. Initial exploration and generation expose query tools only. The public interface does not enable temporal-view isolation, so time language in a question does not hide later records server-side. The training and test split for Experiment 2 is separate from the generation-side patient split. Training uses 20 official tasks and, in the mixed setting, 20 verified synthetic tasks; both settings share 50 official test tasks. The frozen manifest records selected tasks, patients, splits, file hashes, and generation sources.

C Generation Methods and Budgets

TaskCraft maps patient records and read-only FHIR tools onto atomic extraction, verification, and depth/width expansion; both branches start from

atomic tasks, share a generation budget, and record acceptance separately from final filtering. Workflow runs localize–generate–filter with filtering only. SingleShot emits at most two tasks in one pass. SelfRefine drafts, self-critiques, and revises within a patient batch (up to three critiques and two revisions; early stop on pass). EnvExplore obtains environment guidance once, then generates under a fixed initial strategy. MedEvolve uses dual exploration and the full strategy after two rounds of cross-patient feedback-driven updates. All methods share the output schema and external evaluation; internal calls and filters are logged separately.

Matching task counts are not treated as identical token budgets. Exploration and updates count toward MedEvolve cost; shared evaluation is reported separately; reused historical calls are deduplicated from original logs.

D Scoring, Aggregation, and Statistics

D.1 Quality and Reasoning Depth

Quality uses four 1–5 dimensions scored independently: factual accuracy (FA), clinical validity (CV), evidence sufficiency (ES), and question clarity (QC). Rough anchors: 5 = fully accurate/supported/clear; 3 = mostly correct with material gaps; 1 = severe error, unsupported inference, or unusable question. Reasoning depth is labeled from task demand and evidence structure: (1) factual retrieval, (2) single-step judgment, (3) cross-record integration, (4) clinical decision-making; levels 3–4 count as higher-order (HO). Category and depth are recorded separately; treatment-planning items are not auto-labeled as level 4.

The judges are GPT-4o-mini and GPT-5.4-nano. For task i and dimension d, two judges yield $s_{i,d}^{(1)}, s_{i,d}^{(2)}$; we take

$$s_{i,d} = \left\lfloor \operatorname{median}\left(s_{i,d}^{(1)}, s_{i,d}^{(2)}\right) \right\rfloor .$$

Depth uses the same aggregation. The protocol requires two valid judge outputs; a single success does not impute a missing score.

D.2 Evidence Structure and Rate Denominators

Deduplicated cited sources S_i give evidence count $|S_i|$; resource-type count uses recognizable type prefixes. Parseable timestamps T_i define span $\Delta_i = (\max T_i - \min T_i)$ in days (missing if none; zero if one). Span summaries use non-missing values only.

Higher-order task count includes only verified tasks with an aggregate difficulty level of at least 3. Higher-order yield divides this count by the fixed requested-task count. Mean level divides the sum of aggregate levels by the number of verified tasks with difficulty scores. Verification pass rate divides the verified-task count by the actual generated-task count. Experiment 2 divides the number of fully successful test tasks by all 50 planned test tasks.

Generation failures and invalid tasks remain in yield accounting and are not removed from the denominator. Scoring failures in Experiment 2 are reported as missing coverage, rather than relabeled as success or failure. Zero denominators are reported as missing.

D.3 Uncertainty

The higher-order task comparisons in Experiment 1 do not include interval estimates. Experiment 2 uses one seed and 50 planned test tasks.

E Experience Learning and Official Evaluation

The two training settings maintain separate configurations and experience stores and start from the same initial state. Reflexion's cross-task experience distillation is adapted to the medical setting: it stores applicability conditions, operational steps, and lessons from errors, without patient identifiers, individual clinical values, or per-task answers. Each new task retrieves at most three relevant experience entries under frozen rules. Test-time scoring does not trigger reflection or retries.

The execution model is Qwen3.8-27B with thinking disabled, a model-call budget of 100, and seed 42. Both training settings share the same 20 official tasks. The mixed setting adds 20 verified MedEvolve-generated tasks, with 10 generated by GPT-4o-mini and 10 by GPT-5.4-nano.

PhysicianBench uses official tasks and checkpoint-based scoring, with actual file deliverables, execution traces, and the task's FHIR instance. Missing official assets or scoring errors are retained as unavailable results; shared LLM judgments do not replace native scores. Reference answers are used only for scoring and are withheld from execution and learning.

Test scoring is isolated from task solving, and the medical instance used for execution is retained until scoring finishes. Each task uses a copy of the frozen experience. Reference answers, scoring

Framework	Persisted experience	Training updates	Test state
OpenClaw	Skills and an isolated state directory	Native Skill Workshop; whether it is triggered is recorded	Frozen skills; learning and background maintenance disabled
ACE	Playbook	Reflector and Curator updates	Generator only
Reflexion	Reflections and transferable experience	Failure reflection and a limited number of retries	Actor only, reading stored experience

Table 4: Experience representations and update mechanisms of the three frameworks.

results, and full conversations are not written back to the final experience store. Experience hashes are checked before and after testing, and both the planned task count and scoring coverage are reported.

F Generation and Verification Counts

† EnvExplore's PhysicianBench counts come from a frozen earlier batch, rather than the latest run that produced the 21 higher-order tasks in Table 1. These counts cannot serve as the denominator for those 21 tasks.

G Illustrative Cases

Author-verified illustrative cases are pending and are not included in this draft.—

Writing Assistance

Manuscript organization and evidence-status checks were assisted by the scientific-writing skill from Scientific Agent Skills (Kassis et al., 2026). The authors remain responsible for verifying the literature and scientific claims.

	PhysicianBench		MedAgentBench	
Method	Generated / verified	Pass rate	Generated / verified	Pass rate
MedEvolve	157 / 70	44.6%	146 / 104	71.2%
TaskCraft	128 / 125	97.7%	111 / 108	97.3%
Workflow	32 / 10	31.3%	11 / 3	27.3%
SingleShot	162 / 122	75.3%	150 / 127	84.7%
SelfRefine	12 / 10	83.3%	7 / 6	85.7%
EnvExplore	149 / 107†	71.8%†	141 / 108	76.6%

Table 5: Generation and verification counts with GPT-5.4-nano from frozen earlier batches. Entries give actual generated and verified task counts; pass rate divides verified by generated tasks.