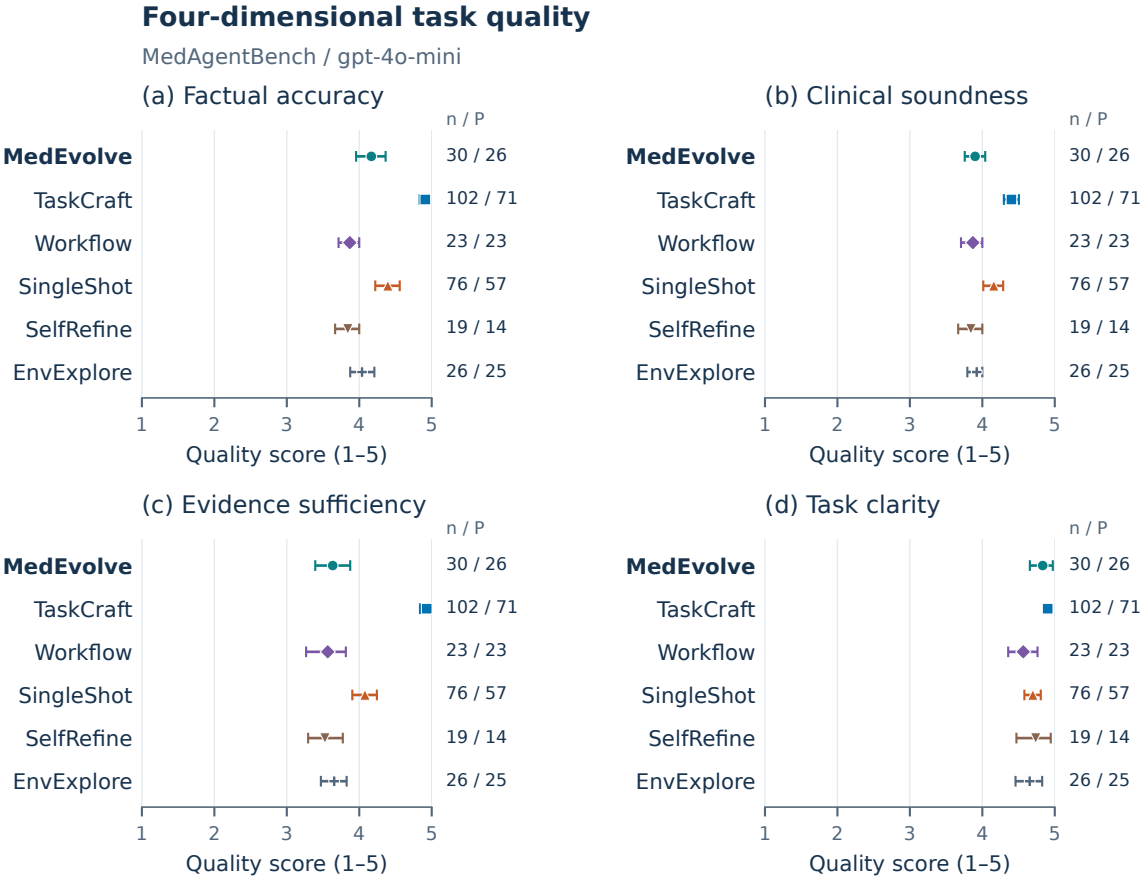
Four-dimensional task quality
MedAgentBench / gpt-4o-mini
(a) Factual accuracy
(b) Clinical soundness
(c) Evidence sufficiency
(d) Task clarity
n / P
MedEvolve
TaskCraft
Workflow
SingleShot
SelfRefine
EnvExplore
30 / 26
102 / 71
23 / 23
76 / 57
19 / 14
26 / 25
1
2
3
4
5
Quality score (1–5)
n: scored tasks P: patients Intervals: patient-cluster 95% bootstrap CI