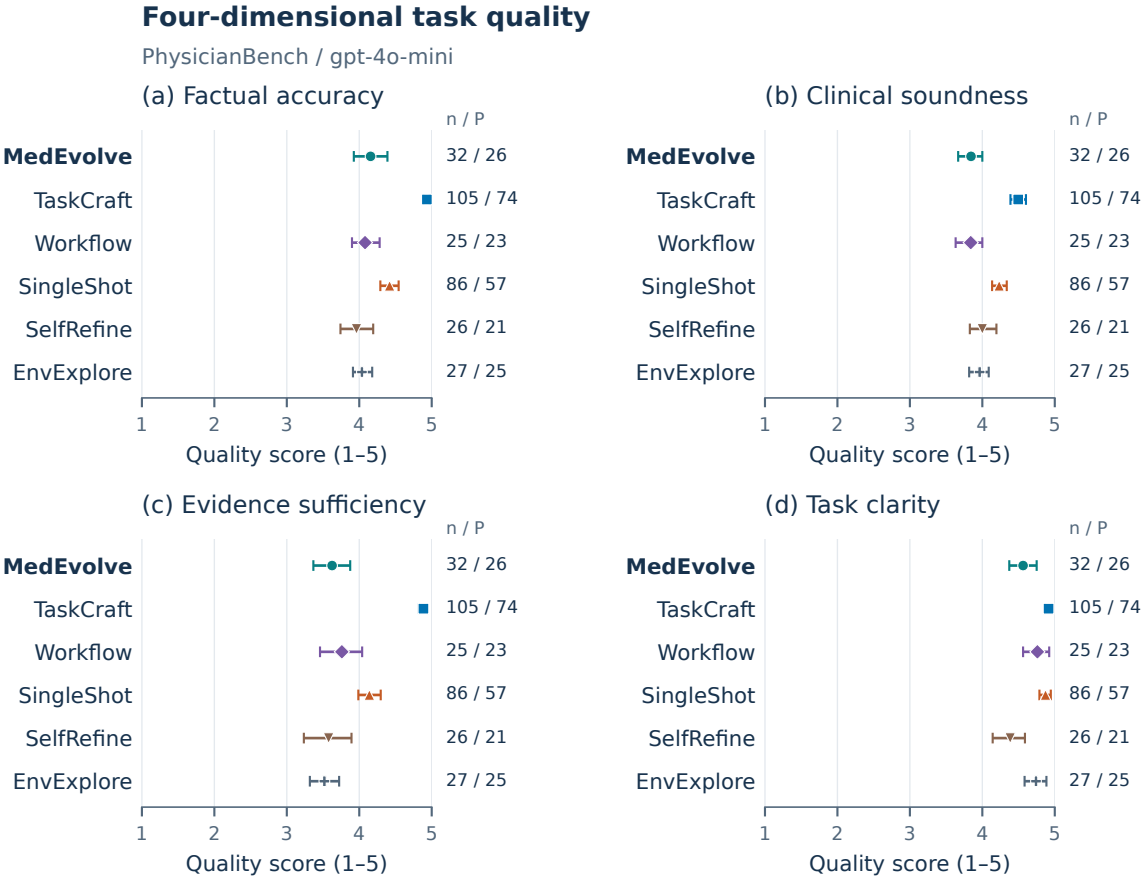
Four-dimensional task quality
PhysicianBench / gpt-4o-mini
(a) Factual accuracy
(b) Clinical soundness
(c) Evidence sufficiency
(d) Task clarity
n / P
MedEvolve
TaskCraft
Workflow
SingleShot
SelfRefine
EnvExplore
32 / 26
105 / 74
25 / 23
86 / 57
26 / 21
27 / 25
1
2
3
4
5
Quality score (1–5)
n: scored tasks P: patients Intervals: patient-cluster 95% bootstrap CI